

YOWON AI

Autonomous AI Jury Platform

OpenHands
Open Source Project

93	ACCEPT	LOW
Overall Score	Deployment Decision	Risk Level

Generated: July 10, 2026 12:47 UTC

Executive Summary

Overview

The overall score is 93, indicating an ACCEPT verdict with a LOW risk level. The agent scores are: technical (86), security (71), scalability (85), innovation (71), and impact (74).

Evaluation Context

Item	Value
Selected Project Type	Open Source Project
AI Detected Type	Open Source Project (59% confidence)
Scoring Rubric Used	Open Source Project
Evaluation Standard	Documentation, community readiness, code quality, maintainability, and contributor experience
Score Band	Exceptional
Confidence	91/100
Evidence Quality	Exceptional
Repository Completeness	100/100


Repository Statistics

Item	Value
Total Files	300
Code Files	200
Documentation Files	28
Presentation Files	0
Test Files	5
Configuration Files	47
Deployment Files	40
Data Files	0
Source Modules	202
Meaningful Files	279
Repository Completeness Score	100

Score Table


Category	Score	Grade
Forge	86/100	A
Sentinel	71/100	B
Scale	85/100	A
Visionary	71/100	B
Guardian	74/100	B

Category Score Chart



Radar Chart

Radar profile of calibrated specialist scores



Risk Visualization

Risk Area	Signal	Risk Level
Security Risk	LOW	
Evidence Gaps	1	
Deployment Readiness	93/100	

Strengths

- Strong community engagement as evidenced by star count
- Good documentation presence
- Self-hosted developer control center

Weaknesses

- Static Analysis Health is at 0/100
- Testing Health is at 55.0/100

Positive Factors

- Source code detected
- Documentation present
- Automated testing detected
- Machine learning implementation
- Dataset artifacts detected
- Backend architecture present
- Deployment files present
- Modular architecture detected

Missing Evidence

- No additional missing evidence detected

Deployment Roadmap

1. Improve Static Analysis Health to 100/100
2. Enhance Testing Health to 100/100

Architecture Summary

Detected frontend, backend, database, ml pipeline, api layer, deployment layer, authentication, integrations, agent systems

Detected Technologies

- FastAPI
- React
- SQLAlchemy

Detected Algorithms

- Agent System
- RAG

Evidence Found

- REST API
- Database
- Authentication
- Docker/deployment
- Testing
- Security implementation
- External integrations
- Agent system

Evidence Missing

- Machine learning model
- Custom algorithm
- Queue/background jobs
- Vector database

Project Type Justification

Detected Open Source Project from code/docs/repository signals.

Calibration Explanation

Implementation evidence increased confidence: REST API, Database, Authentication, Docker/deployment, Testing, Security implementation. Missing evidence primarily reduces confidence unless required by project type: Machine learning model, Custom algorithm, Queue/background jobs, Vector database. Community impact contributed a bounded signal (100/100, capped at 15% of final score). 9 calibration adjustment(s) were applied.

Confidence Sources

- Repository completeness: 100/100
- Documentation quality: present
- Evidence coverage: 100/100

- Agent agreement: inferred from calibrated specialist scores
- Dataset artifact detected

Penalties

- technical: Dataset artifacts strengthen technical evidence (+4)
- technical: REST/API implementation detected (+4)
- technical: Database integration detected (+3)
- technical: External integration evidence detected (+3)
- technical: Open-source rubric: substantial implementation and documentation can score highly
- security: Authentication implementation detected (+4)
- scalability: Advanced architecture component detected (+5)
- impact: Open-source rubric: community activity strengthens impact evidence (+5)
- overall: Community impact signal (+15)

Detailed Findings

Forge Analysis

Forge Analysis

Technical Score:

86/100

Strengths:

Strong community engagement as evidenced by star count

Good documentation presence

Use of Docker for containerization

Weaknesses:

Low test coverage (55%)

Limited number of tests (only 5)

No explicit mention of security practices or policies

Risks:

Sentinel Analysis

Sentinel Analysis

Security Score:

71/100

Risk Level:

LOW

Findings:

None evidenced.

Recommendations:

Run a security review and document controls.

Confidence:

90%

Visionary Assessment

Visionary Analysis

Innovation Score:

71/100

Scalability Score:

85/100

Differentiators:

Self-hosted developer control center

Support for ACP-compatible agents

Multi-backend support

Risks:

Moderate risk due to limited deployment files and no data files.

Recommendations:

Guardian Impact Analysis

Guardian Analysis

Impact Score:

74/100

Top Risks:

Potential security vulnerabilities due to untested code

Inadequate error handling leading to service disruptions

Lack of comprehensive documentation for new contributors

Expected Impact:

Configuration errors in Docker setup

Missing test coverage for critical functions

Dependency issues due to lack of automated testing

Top Failure Predictions

1. Configuration errors in Docker setup
2. Missing test coverage for critical functions
3. Dependency issues due to lack of automated testing

Scalability Assessment

Visionary Scalability Assessment

Scalability Score:

85/100

Risks:

Moderate risk due to limited deployment files and no data files.

Recommendations:

Validate capacity assumptions with load tests and deployment evidence.

Cross Examination Results

No major contradictions detected across specialist agents.

Final Verdict

Blocking Issues

- No blocking issues recorded.

Recommended Fixes

1. Implement Docker Compose Multi-Container Orchestration as per the Top Recommendation

Deployment Roadmap

1. Improve Static Analysis Health to 100/100
2. Enhance Testing Health to 100/100

Deployment Decision

Verdict: ACCEPT
Overall Score: 93/100
Risk Level: LOW